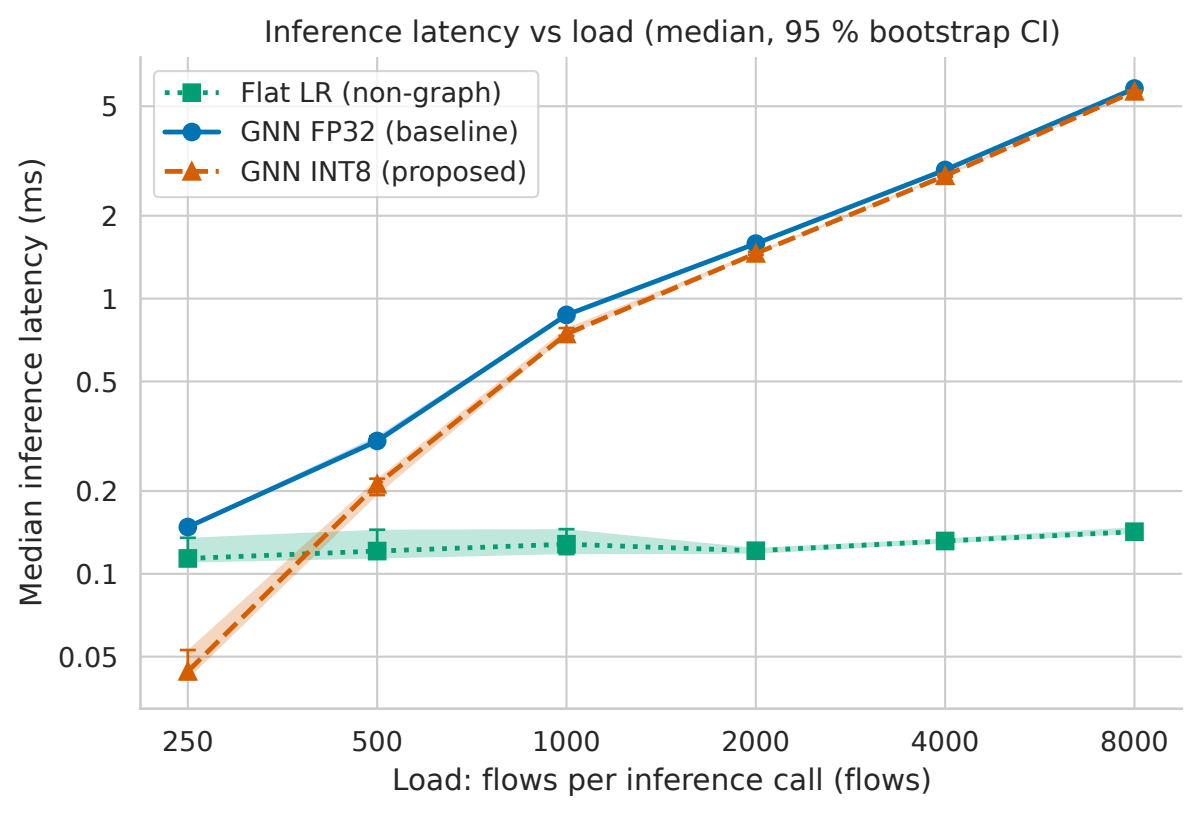
Inference latency vs load (median, 95 % bootstrap CI)
Flat LR (non-graph)
GNN FP32 (baseline)
GNN INT8 (proposed)
Median inference latency (ms)
5
2
1
0.5
0.2
0.1
0.05
250
500
1000
2000
4000
8000
Load: flows per inference call (flows)