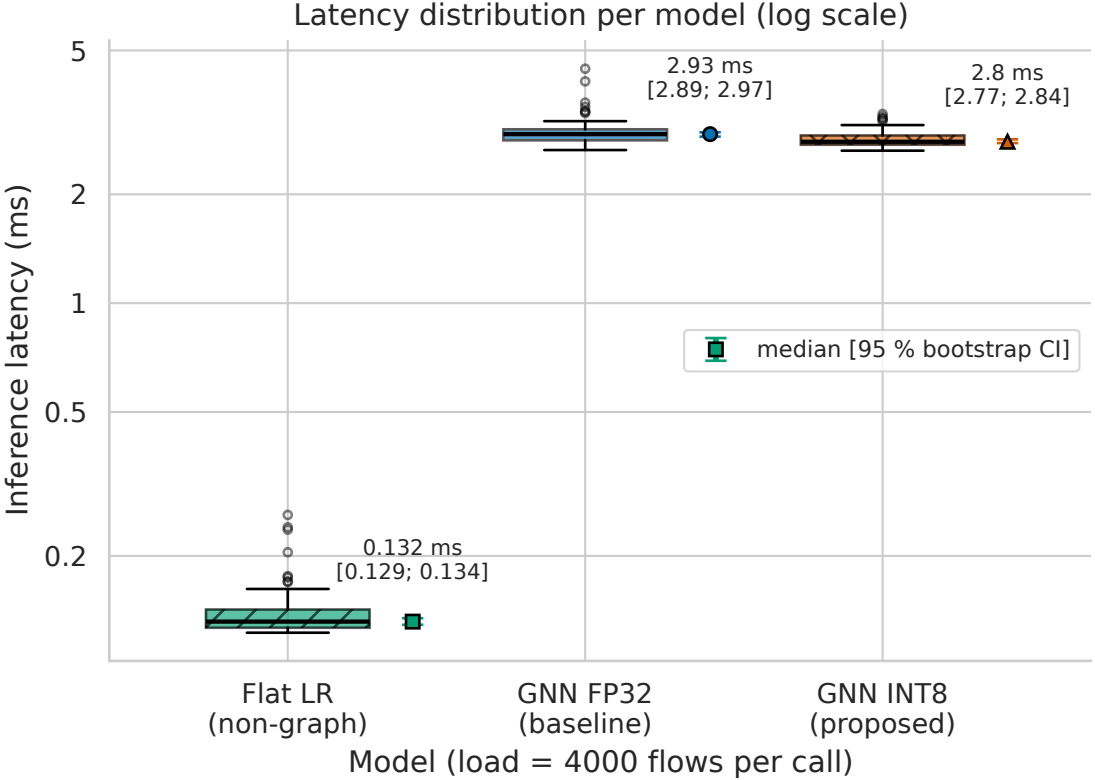
Latency distribution per model (log scale)
Inference latency (ms)
5
2
1
0.5
0.2
2.93 ms
[2.89; 2.97]
2.8 ms
[2.77; 2.84]
0.132 ms
[0.129; 0.134]
median [95 % bootstrap CI]
Flat LR
(non-graph)
GNN FP32
(baseline)
GNN INT8
(proposed)
Model (load = 4000 flows per call)